

Shivam Singh [SHIV]

San Francisco, CA | shiv-pm-ai@umich.edu | +1-650-283-3478 | linkedin.com/in/beastofbayarea

DIRECTOR / PRINCIPAL PRODUCT MANAGER · AI PLATFORM · GENERATIVE & AGENTIC AI

EDUCATION

University of Michigan, Stephen M. Ross School of Business — Master of Business Administration 2022–2024 · Ann Arbor, MI
The London School of Economics and Political Science — Executive Coursework in Economics, Markets & Public Policy 2024 · London, UK
Indian Institute of Technology — Bachelor of Technology, Computer Science 2010–2014

CORE COMPETENCIES

Product Management · AI Platform & Product Strategy · Generative & Agentic AI · LLMs, RAG & Vector Search · ML Systems & MLOps · Product Discovery & User Research · Hypothesis-Driven Experimentation & A/B Testing · Roadmapping & Prioritization (RICE) · Working Backwards / PR-FAQ · 0→1 & Scaling · Activation / Retention / Engagement · Product-Led Growth (PLG) · Lifecycle & Funnel Optimization · Trust & Safety, Guardrails & Evals · Responsible AI & Compliance · Localization (i18n / l10n) · Platform & API Product · Buy-vs-Build, M&A & Partnerships · Telemetry, Analytics & SQL · Stakeholder Management · OKRs / KPIs · GTM

PROFESSIONAL EXPERIENCE

AWS 2024–Present | San Francisco, CA, USA

Principal Product Manager, Technical — Generative AI

Agentic GenAI Shopping Assistant (Consumer, 100M+ Users)

- Owned product strategy and roadmap for an **agentic GenAI shopping assistant**, running **customer discovery** and an **HCI / human-factors redesign**, hardening the **RAG (retrieval-augmented generation)** pipeline, and adding **human-in-the-loop guardrails, evals, and trust-and-safety** controls for high-risk responses.
- Drove **adoption to 25%**, brand favorability **55% → 70%**, **hallucination rate down to 2.8%**, and session depth **+15%** through tight **experimentation, prompt and retrieval tuning, and telemetry-driven iteration**.
- Authored the **PR-FAQ / Working-Backwards** narrative and a **\$100M+ GMV sensitivity model** that secured **\$5M in incremental funding** and executive buy-in for the agentic roadmap.
- Shipped **Japan localization (i18n / l10n)** — SentencePiece tokenization and Keigo honorific handling — as the reusable internationalization template for global market expansion.

GenAI Platform M&A — Buy-vs-Build (Agentic Operating System)

- Led the **buy-vs-build** evaluation and an **\$85M acquisition** of an agentic-AI operating system, defining the **platform make-vs-buy roadmap**, technical due diligence, and the engineering integration plan.
- Structured **ROFR/ROFN** protections and drove a post-deal **unit-economics** turnaround, lifting **gross margin 55% → 81%** while consolidating overlapping model and tooling spend.

Rakuten 2023 | Tokyo, Japan

Senior Product Manager, AI Platform & Merchant Solutions

Enterprise GenAI Platform — 70+ Business Units, 50K+ Merchants

- Built a horizontal **GenAI platform spanning 70+ business units and 50,000+ merchants** on a “buy intelligence, build context” thesis, integrating a **foundation-model partnership** behind a **PII redaction gateway** and **hybrid vector + keyword (semantic) search** over a unified knowledge graph.
- Delivered **+15% merchant conversion uplift**, cut listing time **45 → 12 min**, reached **40% platform adoption in 6 months**, and reduced network energy use **25%** via inference optimization and caching.
- Exercised hard **product judgment** — killed an underperforming chatbot UI after usage analysis and re-scoped toward embedded merchant workflows that moved the conversion metric.

Microsoft 2020–2022 | Redmond, WA, USA

Senior Product Manager, Lifecycle & Growth

Developer + Business Lifecycle for an LLM Platform (Product-Led Growth)

- Owned the **developer + business lifecycle** for an LLM platform end-to-end — onboarding, activation, expansion, and win-back — instrumented with **cohort telemetry, funnel analytics, and a continuous experimentation backlog**.
- Lifted activation conversion **+262%**, grew **net revenue retention (NRR) 102% → 128%**, cut churn **14% → 6.5%**, and improved win-back **4% → 22%** through **PLG** loops and time-to-first-value optimization.
- Translated lifecycle and usage signals into the **product roadmap**, prioritizing (via **RICE**) the onboarding and first-run-experience work that moved activation most, in lockstep with engineering, design, and developer-relations.
- Partnered with data science to build the **retention and propensity models** that targeted lifecycle interventions and forecast expansion revenue.

The D. E. Shaw Group 2016–2020 | New York, NY, USA & Hyderabad, India

Vice President, Product — Investment Infrastructure

Intelligent Document Processing (IDP) Platform & Deal Operations

- Led product for an **Intelligent Document Processing (IDP)** platform that replaced manual data entry with **semantic extraction** and a **confidence-scored, human-in-the-loop** pipeline (auto-pass >98% confidence; route <98% for review, with every correction fed back as **training data / active learning**).
- Collapsed earnings-report ingestion from hours to seconds at **~70% cost savings**, then built a unified deal-data repository in a **72-hour sprint** that cut consolidation time **48h → 6h (87.5%)** and helped close a **\$25M Series B** a week ahead of competitors.
- Ran the platform as an **internal product** — treating quant, trading, and legal teams as customers, defining SLAs and adoption metrics, and prioritizing the backlog against deal-velocity impact.

McKinsey & Company 2014–2016 | London, UK

Associate, Digital & Product (Fintech)

Neobank Launch — Alternative-Data Credit Scoring (0→1)

- Drove **0→1 product launch** of a SE-Asia neobank built on **alternative telco-data credit scoring** (ML across 400+ features), a **regulatory sandbox**, and **eKYC liveness detection**, embedded into a super-app for distribution.
- Scaled **0 → 100K active users in 90 days**, lifted loan approval rate **10% → 40%**, and held a **2.4% non-performing-loan (NPL)** rate at prime-bank parity while keeping **CAC under \$10**.
- Owned the **discovery-to-launch** arc — market sizing, MVP scoping, model risk-and-fairness review, and go-to-market — coordinating product, data science, risk, and compliance.